

HW3

October 25, 2022

```
[1]: import numpy as np
import matplotlib.pyplot as plt
import pandas as pd
from sklearn.datasets import load_breast_cancer
from sklearn import datasets
from sklearn import metrics
from sklearn.naive_bayes import GaussianNB
import seaborn as sns
from sklearn.linear_model import LogisticRegression
from sklearn.model_selection import train_test_split
```

```
[2]: breast = load_breast_cancer()
```

```
[3]: breast_data = breast.data
breast_data.shape
```

```
[3]: (569, 30)
```

```
[4]: breast_input = pd.DataFrame(breast_data)
breast_input.head()
```

```
[4]:
```

	0	1	2	3	4	5	6	7	8	\		
0	17.99	10.38	122.80	1001.0	0.11840	0.27760	0.3001	0.14710	0.2419			
1	20.57	17.77	132.90	1326.0	0.08474	0.07864	0.0869	0.07017	0.1812			
2	19.69	21.25	130.00	1203.0	0.10960	0.15990	0.1974	0.12790	0.2069			
3	11.42	20.38	77.58	386.1	0.14250	0.28390	0.2414	0.10520	0.2597			
4	20.29	14.34	135.10	1297.0	0.10030	0.13280	0.1980	0.10430	0.1809			
		9	...	20	21	22	23	24	25	26	27	\
0	0.07871	...	25.38	17.33	184.60	2019.0	0.1622	0.6656	0.7119	0.2654		
1	0.05667	...	24.99	23.41	158.80	1956.0	0.1238	0.1866	0.2416	0.1860		
2	0.05999	...	23.57	25.53	152.50	1709.0	0.1444	0.4245	0.4504	0.2430		
3	0.09744	...	14.91	26.50	98.87	567.7	0.2098	0.8663	0.6869	0.2575		
4	0.05883	...	22.54	16.67	152.20	1575.0	0.1374	0.2050	0.4000	0.1625		
		28		29								
0	0.4601	0.11890										
1	0.2750	0.08902										

```

2  0.3613  0.08758
3  0.6638  0.17300
4  0.2364  0.07678

```

[5 rows x 30 columns]

```
[5]: breast_labels = breast.target
breast_labels.shape
```

```
[5]: (569,)
```

```
[6]: labels = np.reshape(breast_labels,(569,1))
final_breast_data = np.concatenate([breast_data,labels],axis=1)
final_breast_data.shape
```

```
[6]: (569, 31)
```

```
[7]: breast_dataset = pd.DataFrame(final_breast_data)
features = breast.feature_names
features
```

```
[7]: array(['mean radius', 'mean texture', 'mean perimeter', 'mean area',
        'mean smoothness', 'mean compactness', 'mean concavity',
        'mean concave points', 'mean symmetry', 'mean fractal dimension',
        'radius error', 'texture error', 'perimeter error', 'area error',
        'smoothness error', 'compactness error', 'concavity error',
        'concave points error', 'symmetry error',
        'fractal dimension error', 'worst radius', 'worst texture',
        'worst perimeter', 'worst area', 'worst smoothness',
        'worst compactness', 'worst concavity', 'worst concave points',
        'worst symmetry', 'worst fractal dimension'], dtype='<U23')
```

```
[8]: features_labels = np.append(features,'Type')
breast_dataset.columns = features_labels
breast_dataset.head()
```

```
[8]:
```

	mean radius	mean texture	mean perimeter	mean area	mean smoothness	\
0	17.99	10.38	122.80	1001.0	0.11840	
1	20.57	17.77	132.90	1326.0	0.08474	
2	19.69	21.25	130.00	1203.0	0.10960	
3	11.42	20.38	77.58	386.1	0.14250	
4	20.29	14.34	135.10	1297.0	0.10030	

	mean compactness	mean concavity	mean concave points	mean symmetry	\
0	0.27760	0.3001	0.14710	0.2419	
1	0.07864	0.0869	0.07017	0.1812	
2	0.15990	0.1974	0.12790	0.2069	
3	0.28390	0.2414	0.10520	0.2597	

4	0.13280	0.1980	0.10430	0.1809
---	---------	--------	---------	--------

	mean fractal dimension	...	worst texture	worst perimeter	worst area	\
0	0.07871	...	17.33	184.60	2019.0	
1	0.05667	...	23.41	158.80	1956.0	
2	0.05999	...	25.53	152.50	1709.0	
3	0.09744	...	26.50	98.87	567.7	
4	0.05883	...	16.67	152.20	1575.0	

	worst smoothness	worst compactness	worst concavity	worst concave points	\
0	0.1622	0.6656	0.7119	0.2654	
1	0.1238	0.1866	0.2416	0.1860	
2	0.1444	0.4245	0.4504	0.2430	
3	0.2098	0.8663	0.6869	0.2575	
4	0.1374	0.2050	0.4000	0.1625	

	worst symmetry	worst fractal dimension	Type
0	0.4601	0.11890	0.0
1	0.2750	0.08902	0.0
2	0.3613	0.08758	0.0
3	0.6638	0.17300	0.0
4	0.2364	0.07678	0.0

[5 rows x 31 columns]

```
[9]: breast_dataset['Type'].replace(0, 'Benign',inplace=True)
breast_dataset['Type'].replace(1, 'Malignant',inplace=True)
```

```
[10]: breast_dataset.tail()
```

```
[10]:
```

	mean radius	mean texture	mean perimeter	mean area	mean smoothness	\
564	21.56	22.39	142.00	1479.0	0.11100	
565	20.13	28.25	131.20	1261.0	0.09780	
566	16.60	28.08	108.30	858.1	0.08455	
567	20.60	29.33	140.10	1265.0	0.11780	
568	7.76	24.54	47.92	181.0	0.05263	

	mean compactness	mean concavity	mean concave points	mean symmetry	\
564	0.11590	0.24390	0.13890	0.1726	
565	0.10340	0.14400	0.09791	0.1752	
566	0.10230	0.09251	0.05302	0.1590	
567	0.27700	0.35140	0.15200	0.2397	
568	0.04362	0.00000	0.00000	0.1587	

	mean fractal dimension	...	worst texture	worst perimeter	worst area	\
564	0.05623	...	26.40	166.10	2027.0	
565	0.05533	...	38.25	155.00	1731.0	

566	0.05648	...	34.12	126.70	1124.0
567	0.07016	...	39.42	184.60	1821.0
568	0.05884	...	30.37	59.16	268.6

	worst smoothness	worst compactness	worst concavity \
564	0.14100	0.21130	0.4107
565	0.11660	0.19220	0.3215
566	0.11390	0.30940	0.3403
567	0.16500	0.86810	0.9387
568	0.08996	0.06444	0.0000

	worst concave points	worst symmetry	worst fractal dimension	Type
564	0.2216	0.2060	0.07115	Benign
565	0.1628	0.2572	0.06637	Benign
566	0.1418	0.2218	0.07820	Benign
567	0.2650	0.4087	0.12400	Benign
568	0.0000	0.2871	0.07039	Malignant

[5 rows x 31 columns]

```
[11]: # 80% Training split
from sklearn.preprocessing import StandardScaler
y = breast_dataset.loc[:, ['Type']].values
x = breast_dataset.loc[:, features].values
x = StandardScaler().fit_transform(x)
X_train, X_test, Y_train, Y_test = train_test_split(x, y, train_size=0.8,
↳test_size=0.2)
y[:10]
```

```
[11]: array(['Benign',
            'Benign',
            'Benign',
            'Benign',
            'Benign',
            'Benign',
            'Benign',
            'Benign',
            'Benign',
            'Benign'], dtype=object)
```

```
[12]: #Problem 1
```

```
[13]: model = GaussianNB()
```

```
[14]: model.fit(X_train, Y_train.ravel())
print(model)
```

GaussianNB()

```
[15]: expected = Y_test
      predicted = model.predict(X_test)
```

```
[16]: print(metrics.classification_report(expected, predicted))
      print(metrics.confusion_matrix(expected, predicted))
```

	precision	recall	f1-score	support
Benign	0.88	0.88	0.88	33
Malignant	0.95	0.95	0.95	81
accuracy			0.93	114
macro avg	0.91	0.91	0.91	114
weighted avg	0.93	0.93	0.93	114

```
[[29  4]
 [ 4 77]]
```

```
[17]: # Compare with Logistic Regression
      classifier = LogisticRegression(penalty='none', random_state=0)
```

```
[18]: classifier.fit(X_train, Y_train.ravel())
```

```
[18]: LogisticRegression(penalty='none', random_state=0)
```

```
[19]: Y_pred = classifier.predict(X_test)
```

```
[20]: print(metrics.classification_report(expected, Y_pred))
      print(metrics.confusion_matrix(expected, Y_pred))
```

	precision	recall	f1-score	support
Benign	0.91	0.94	0.93	33
Malignant	0.97	0.96	0.97	81
accuracy			0.96	114
macro avg	0.94	0.95	0.95	114
weighted avg	0.96	0.96	0.96	114

```
[[31  2]
 [ 3 78]]
```

```
[21]: # Problem 2


      #Decide the number of PCA components
      from sklearn.decomposition import PCA
      pca = PCA(random_state=88)
      pca.fit(X_train)
```

```

explained_variance = np.cumsum(pca.explained_variance_ratio_)
plt.vlines(x=9, ymax=1, ymin=0, colors='r', linestyle="--")
plt.hlines(y=0.95, xmax=120, xmin=0, colors="g", linestyle="--")
plt.plot(explained_variance)
#principalComponents = pca.fit_transform(x)
#principalDf = pd.DataFrame(data=principalComponents
#                           , columns=['principal component 1', 'principal_
↪component 2'])

```

[21]: [<matplotlib.lines.Line2D at 0x2096357b580>]



```


[22]: # Train PCA model
pca_final = PCA(0.95)
df_train_pca = pca_final.fit_transform(X_train)

```

```

[23]: # Correlations between components
corr_mat = np.corrcoef(df_train_pca.transpose())
plt.figure(figsize=[15,8])
sns.heatmap(corr_mat)
plt.show()

```



```
[24]: # Apply PCA model to test data
df_test_pca = pca_final.transform(X_test)
# Train Logistic Regression
LR = LogisticRegression()
LR_model = LR.fit(df_train_pca, Y_train.ravel())
```

```
[25]: # Model Evaluation
Y_pred_pca = LR_model.predict(df_test_pca)
cnf_matrix = metrics.confusion_matrix(Y_test, Y_pred_pca)
cnf_matrix
```

```
[25]: array([[31,  2],
        [ 2, 79]], dtype=int64)
```


```
[26]: # Evaluate accuracy, precision, and recall
print("Accuracy:", metrics.accuracy_score(Y_test, Y_pred_pca))
print("Precision:", metrics.precision_score(Y_test, Y_pred_pca,
    pos_label='Benign'))
print("Recall:", metrics.recall_score(Y_test, Y_pred_pca, pos_label='Benign'))
Accuracy = []
Precision = []
Recall = []
k = range(1, len(features)+1)
```

```
Accuracy: 0.9649122807017544
Precision: 0.9393939393939394
```

Recall: 0.9393939393939394

```
[27]: # Plot Model Performance vs # of Principal Components
for i in k:
    pca_final = PCA(n_components=i)
    df_train_pca = pca_final.fit_transform(X_train)
    df_test_pca = pca_final.transform(X_test)
    LR = LogisticRegression()
    LR_model = LR.fit(df_train_pca, Y_train.ravel())
    Y_pred_pca = LR_model.predict(df_test_pca)
    #print("K = ", i)
    #print("Accuracy:", metrics.accuracy_score(Y_test, Y_pred_pca))
    #print("Precision:", metrics.precision_score(Y_test, Y_pred_pca,
    pos_label='Benign'))
    #print("Recall:", metrics.recall_score(Y_test, Y_pred_pca,
    pos_label='Benign'))
    Accuracy.append(metrics.accuracy_score(Y_test, Y_pred_pca))
    Precision.append(metrics.precision_score(Y_test, Y_pred_pca,
    pos_label='Benign'))
    Recall.append(metrics.recall_score(Y_test, Y_pred_pca, pos_label='Benign'))

[28]: plt.plot(k, Accuracy, 'b', label="Accuracy")
plt.plot(k, Precision, 'g', label="Precision")
plt.plot(k, Recall, 'r', label="Recall")
plt.xlabel("# of Principal Components")
plt.ylabel("Model Performance")
plt.legend()
plt.show()
```

[29]: # Problem 3

```
[30]: model = GaussianNB()
model.fit(df_train_pca, Y_train.ravel())
k = range(1, len(features)+1)
Accuracy_NB = []
Precision_NB = []
Recall_NB = []
```

```
[31]: for i in k:
    pca_final = PCA(n_components=i)
    df_train_pca = pca_final.fit_transform(X_train)
    df_test_pca = pca_final.transform(X_test)
    model = GaussianNB()
    model.fit(df_train_pca, Y_train.ravel())
    Y_pred_pca = model.predict(df_test_pca)
    #DEBUG print("K = ", i)
    #DEBUG print("Accuracy:", metrics.accuracy_score(Y_test, Y_pred_pca))
    #DEBUG print("Precision:", metrics.precision_score(Y_test, Y_pred_pca, ↵
↵pos_label='Benign'))
    #DEBUG print("Recall:", metrics.recall_score(Y_test, Y_pred_pca, ↵
↵pos_label='Benign'))
    Accuracy_NB.append(metrics.accuracy_score(Y_test, Y_pred_pca))
```

```

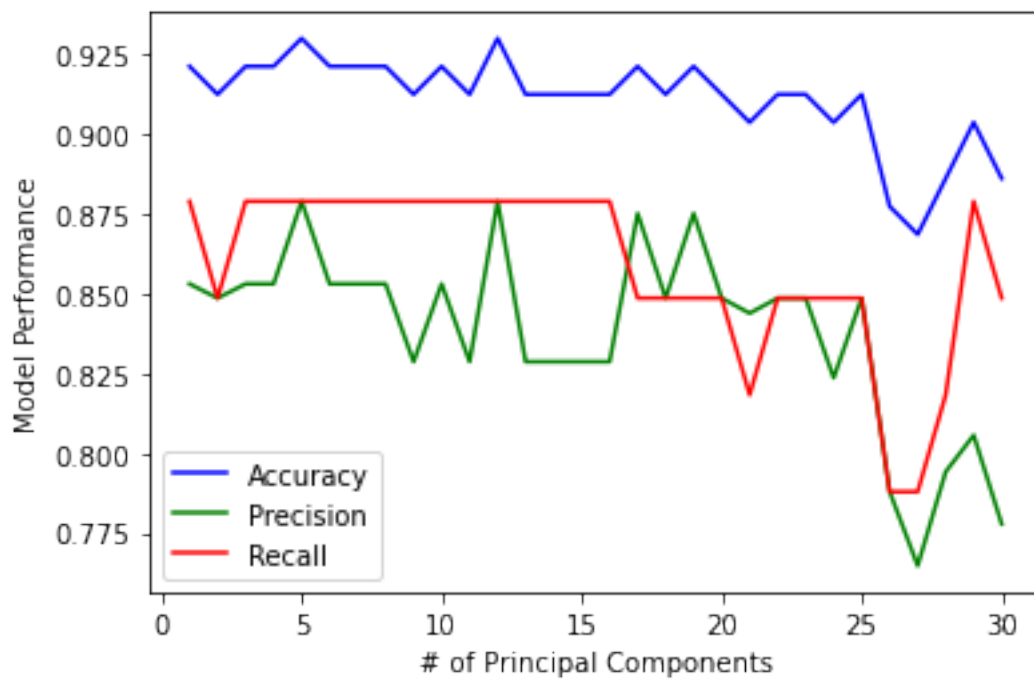
    Precision_NB.append(metrics.precision_score(Y_test, Y_pred_pca,
↪pos_label='Benign'))
    Recall_NB.append(metrics.recall_score(Y_test, Y_pred_pca,
↪pos_label='Benign'))

```

```

[32]: plt.plot(k, Accuracy_NB, 'b', label="Accuracy")
plt.plot(k, Precision_NB, 'g', label="Precision")
plt.plot(k, Recall_NB, 'r', label="Recall")
plt.xlabel("# of Principal Components")
plt.ylabel("Model Performance")
plt.legend()
plt.show()

```



```

[ ]:

```